

---

## Fractal Workspace

### Visual-first, node-based idea workspace for human + LLM collaboration

**Tagline:** Think in branches. Keep the truth. Scale memory beyond the context window.

---

### Executive Summary (Elevator Pitch)

Fractal Workspace introduces a new interaction primitive for LLMs: a **visual, node-based workspace** where each node is a bounded chat instance with its own frozen context, a generated name, and a structured contribution to a lineage knowledge graph.

This approach addresses three persistent problems in LLM-assisted thinking:

1. Context pollution when exploring multiple ideas in parallel
2. Short-lived, brittle context memory and hallucinations as conversations grow
3. Practical performance limits on local or smaller LLMs (memory and speed)

By combining **branching, knowledge-graph summaries, and retrieval-augmented generation (RAG)**, Fractal Workspace enables reliable, scalable, and explainable ideation workflows — even on constrained local hardware.

---

### The Problem (Why This Matters)

As large language models have become more capable, their limitations during *extended reasoning and ideation* have become increasingly apparent.

Humans naturally think in **branches**. When exploring ideas, we maintain parallel lines of thought and later merge or discard them. Traditional chat interfaces are strictly linear, forcing unrelated threads into a single context and degrading output quality over time.

LLMs also suffer from **context degradation**. As the context window fills, models become more prone to hallucination and factual drift. This is a well-documented phenomenon in neural text generation and abstractive summarization research.

Finally, **local inference constraints** remain a practical concern. Quantization techniques such as 4-bit and QLoRA make local deployment feasible, but memory bandwidth and latency still demand careful control over context size.

In short: current chat interfaces do not align with how humans think, nor with how LLMs remain reliable.

---

## Reframing the Problem: Fractal Workspace

Fractal Workspace replaces the single linear chat with a **directed acyclic graph (DAG) of semantic states**.

Each **node** represents a bounded chat instance with:

- A frozen conversation log
- A generated short name
- A derived summary
- A local knowledge-graph delta (entities and relations introduced)
- A lifecycle state (active, frozen, deleted)

**Edges** represent semantic lineage rather than raw message flow.

For any LLM call, the effective context is constructed as:

System prompt

- Lineage summary
- Lineage knowledge graph
- Node summary
- Node graph delta
- Last-N node-local messages

Merges are explicit operations. A merge produces a new summary and graph delta, which are appended as derived artifacts to the target node. Forgetting becomes intentional, traceable, and auditable.

---

## Implemented Features

- Visual infinite-canvas UI with first-class nodes
- Branching, drag-to-merge, copy, delete, freeze operations
- Event-sourced backend (append-only action log as source of truth)
- Node-level summarization with provenance tracking

- Knowledge-graph construction (entity-relation deltas per node)
  - Retrieval-augmented generation (RAG) using embeddings
  - Auto-splitting agent for overloaded nodes (experimental)
  - Local-first LLM orchestration using Ollama
  - Quantization-aware model selection per role
  - Traceable deletion with optional downstream recompute
- 

### Example 1: Branching Prevents Context Pollution

**Scenario:** Designing a neonatal jaundice research pipeline.

- Center node: “Neonatal Jaundice Project” (goals and constraints)
- Branch A: “Data Pipeline” (data ingestion, bilirubin trends, calibration)
- Branch B: “Model Architecture” (CNN vs XGBoost vs time-series models)

In a linear chat, low-level preprocessing details from Branch A bleed into architectural reasoning in Branch B. In Fractal Workspace, Branch B inherits only the **explicit summary** of the center node. It never sees data-pipeline chatter unless explicitly merged.

When the user decides to merge, the system generates a concise merged summary and graph delta — the only information promoted to global context.

---

### Example 2: Knowledge Graph + Summary Compression

After six sessions exploring multiple models and preprocessing strategies, a linear chat would require thousands of tokens of history.

In Fractal Workspace:

- Each node emits graph deltas such as:
  - Decision: use EfficientNet-based CNN
  - Experiment: augmentation A and B failed
- The summarizer compresses lineage into a short, decision-focused summary
- Provenance is preserved via node IDs

Future LLM calls receive only the distilled meaning plus pointers, not the full raw logs.  
Token usage stays low while auditability remains intact.

---

## **Retrieval-Augmented Generation (RAG)**

RAG enables effectively unlimited long-term memory by storing embeddings of important artifacts and retrieving only the most relevant information at query time.

This approach:

- Reduces hallucination on knowledge-intensive tasks
- Keeps prompt sizes minimal
- Works well with local vector databases such as FAISS or Qdrant

Fractal Workspace integrates RAG as an optional augmentation layer whenever the agent detects missing context.

---

## **Why Hallucination and Forgetting Occur**

- Abstractive models frequently hallucinate when compressing or inferring information
- Larger context windows mitigate but do not eliminate this issue
- Retrieval and structured memory are widely accepted mitigation strategies

These limitations are not theoretical — they are empirically documented across summarization and dialogue tasks.

---

## **Why Bigger Models Alone Are Not the Solution**

Increasing parameter count and context window size is expensive and often impractical, especially for local inference.

Fractal Workspace focuses instead on **memory architecture**:

- Explicit structure
- Provenance tracking
- Controlled context construction

Quantization techniques like QLoRA make large models accessible, but disciplined context management is still essential. This system enforces that discipline by design.

---

## Technical Architecture Overview

Clients (Browser UI)

↔ API (FastAPI)

↔ Event Store (Append-only)

Event Store → Derived Artifacts (Summaries, Graph Deltas)

Embeddings / Vector DB ↔ Retriever

LLM Layer (Chat, Summarizer, Graph Builder, Node Namer via Ollama)

Agents subscribe to events and trigger summarization, retrieval, and auto-splitting.

---

## Demo Flow

1. Create center node and enter project description
  2. Create two branches and demonstrate isolated reasoning
  3. Introduce conflicting decisions and perform a merge
  4. Show generated summary and knowledge-graph delta
  5. Trigger retrieval for missing context
  6. Demonstrate auto-split of an overloaded node
  7. Inspect event log to show full traceability
- 

## Limitations and Tradeoffs

- Manual summarization in Phase-1
  - Downstream recompute on deletion can be expensive
  - RAG quality depends on embedding and chunking strategy
  - Local inference still constrained by GPU memory and concurrency limits
-

## **Repository Structure (Phase-1)**

backend/ — FastAPI event-sourced API  
agent/ — auto-split and orchestration logic  
llm-config/ — Ollama Modelfiles and role configs  
docs/ — slides and documentation

---

## **How to Run**

1. Start Ollama locally and load a supported model
  2. Run backend using provided script
  3. Open API docs and UI
  4. Follow demo script
- 

## **Final Note**

This system aligns:

- How humans think
- How LLMs remain reliable
- How local inference constraints work in practice

Branching, structured memory, retrieval, and provenance are not optional features — they are necessary foundations for serious LLM-assisted reasoning.

Contributors:

Aryan Guptas 2<sup>nd</sup> year mechanical engg(IITR)

Suryansh Singh 2nd year mechanical engg(IITR)